

ETL Project

ETL: Extract, Transform, Load

For this project, I wanted to see if the amount of sleep I got affected my blood glucose values.

Extract:

Two data sources:

1. Dexcom glucose values
2. Fitbit sleep values

Transform:

Transform data and data cleaning.

```
In [1]: # Import dependencies

import pandas as pd
from sqlalchemy import create_engine
```

```
In [2]: # Read the CSV file

csv_file_1 = "Resources/DEXCOM_CGM_08082019_to_11052019.csv"
cgm_data_df = pd.read_csv(csv_file_1)

# Display the data

cgm_data_df.head()
```

Out[2]:

	Index	Timestamp (YYYY-MM-DDThh:mm:ss)	Event Type	Event Subtype	Patient Info	Device Info	Source Device ID	Glucose Value (mg/dL)	Insulin Value (u)	Carb Value (grams)	(l
0	1	2019-08-08T00:01:13	EGV	NaN	NaN	NaN	iPhone G6	140	NaN	NaN	
1	2	2019-08-08T00:06:13	EGV	NaN	NaN	NaN	iPhone G6	138	NaN	NaN	
2	3	2019-08-08T00:11:13	EGV	NaN	NaN	NaN	iPhone G6	136	NaN	NaN	
3	4	2019-08-08T00:16:13	EGV	NaN	NaN	NaN	iPhone G6	134	NaN	NaN	
4	5	2019-08-08T00:21:13	EGV	NaN	NaN	NaN	iPhone G6	131	NaN	NaN	

```
In [3]: # Get only the columns needed from the glucose values CSV

new_cgm_data_df = cgm_data_df[['Timestamp (YYYY-MM-DDThh:mm:ss)', 'Event Type',
, 'Glucose Value (mg/dL)']].copy()

# Display the data

new_cgm_data_df.head()
```

Out[3]:

	Timestamp (YYYY-MM-DDThh:mm:ss)	Event Type	Glucose Value (mg/dL)
0	2019-08-08T00:01:13	EGV	140
1	2019-08-08T00:06:13	EGV	138
2	2019-08-08T00:11:13	EGV	136
3	2019-08-08T00:16:13	EGV	134
4	2019-08-08T00:21:13	EGV	131

```
In [4]: # Do the same for the sleep CSV

csv_file_2 = "Resources/fitbit_sleep_08082019_to_11052019.csv"
sleep_data_df = pd.read_csv(csv_file_2)

# Display the data

sleep_data_df.head()
```

Out[4]:

	Start Time	End Time	Minutes Asleep	Minutes Awake	Number of Awakenings	Time in Bed	Minutes REM Sleep	Minutes Light Sleep	Minutes Deep Sleep
0	2019-11-03 12:18AM	2019-11-03 8:13AM	429	46	30	475	131.0	210.0	88.0
1	2019-11-02 3:47PM	2019-11-02 7:30PM	192	31	19	223	29.0	135.0	28.0
2	2019-11-01 8:12PM	2019-11-02 7:21AM	574	95	40	669	120.0	379.0	75.0
3	2019-11-01 2:08AM	2019-11-01 7:00AM	247	45	16	292	50.0	161.0	36.0
4	2019-10-30 8:35PM	2019-10-31 6:36AM	508	93	2	601	NaN	NaN	NaN

```
In [5]: # Get only the columns needed from the sleep CSV

new_sleep_data_df = sleep_data_df[['End Time', 'Minutes Asleep', 'Minutes Awake', 'Number of Awakenings', 'Time in Bed', 'Minutes REM Sleep', 'Minutes Light Sleep', 'Minutes Deep Sleep']].copy()

# Display the data

new_sleep_data_df.head()
```

Out[5]:

	End Time	Minutes Asleep	Minutes Awake	Number of Awakenings	Time in Bed	Minutes REM Sleep	Minutes Light Sleep	Minutes Deep Sleep
0	2019-11-03 8:13AM	429	46	30	475	131.0	210.0	88.0
1	2019-11-02 7:30PM	192	31	19	223	29.0	135.0	28.0
2	2019-11-02 7:21AM	574	95	40	669	120.0	379.0	75.0
3	2019-11-01 7:00AM	247	45	16	292	50.0	161.0	36.0
4	2019-10-31 6:36AM	508	93	2	601	NaN	NaN	NaN

```
In [6]: # Convert the time stamp to display only the date
# Because I will be combining the data to do analysis later, the date will be the joining column
# The time stamp will be too specific so I need to make sure data displays date only

new_cgm_data_df['Timestamp (YYYY-MM-DDThh:mm:ss)'] = pd.to_datetime(new_cgm_data_df['Timestamp (YYYY-MM-DDThh:mm:ss)']).dt.strftime('%Y-%m-%d')

# Display the data

new_cgm_data_df.head()
```

Out[6]:

	Timestamp (YYYY-MM-DDThh:mm:ss)	Event Type	Glucose Value (mg/dL)
0	2019-08-08	EGV	140
1	2019-08-08	EGV	138
2	2019-08-08	EGV	136
3	2019-08-08	EGV	134
4	2019-08-08	EGV	131

```
In [7]: # Check the amount of data

len(new_cgm_data_df)
```

Out[7]: 25772

```
In [8]: # Drop all Event Type values that are not equal to EGV (estimated glucose value)

new_cgm_data_df = new_cgm_data_df[~(new_cgm_data_df['Event Type'] != 'EGV')]

# Display the data

new_cgm_data_df.head()
```

Out[8]:

	Timestamp (YYYY-MM-DDThh:mm:ss)	Event Type	Glucose Value (mg/dL)
0	2019-08-08	EGV	140
1	2019-08-08	EGV	138
2	2019-08-08	EGV	136
3	2019-08-08	EGV	134
4	2019-08-08	EGV	131

```
In [9]: # Check count to see if rows were dropped

len(new_cgm_data_df)
```

Out[9]: 25449

```
In [10]: # Rename columns for clean data

new_cgm_data_df = new_cgm_data_df.rename(columns={'Timestamp (YYYY-MM-DDThh:mm:ss)': 'date_entry', 'Event Type': 'event_type', 'Glucose Value (mg/dL)': 'glucose_value'})

# Display the data

new_cgm_data_df.head()
```

Out[10]:

	date_entry	event_type	glucose_value
0	2019-08-08	EGV	140
1	2019-08-08	EGV	138
2	2019-08-08	EGV	136
3	2019-08-08	EGV	134
4	2019-08-08	EGV	131

```
In [11]: # Drop the event_type column because I will not need it for further analysis

new_cgm_data_df = new_cgm_data_df.drop(columns=['event_type'])

# Drop NaN values from the table

new_cgm_data_df = new_cgm_data_df.dropna()
```

In [13]: *# Check data types of the values - will need to convert later*

```
new_cgm_data_df.dtypes
```

Out[13]:

date_entry	object
glucose_value	object
dtype:	object

In [14]: *# Get rid of all non-values such as "High" and "Low"*

```
new_cgm_data_df = new_cgm_data_df[(new_cgm_data_df['glucose_value'] != 'High')]  
new_cgm_data_df = new_cgm_data_df[(new_cgm_data_df['glucose_value'] != 'Low')]  
  
# Display the data  
  
new_cgm_data_df.head()
```

Out[14]:

	date_entry	glucose_value
0	2019-08-08	140
1	2019-08-08	138
2	2019-08-08	136
3	2019-08-08	134
4	2019-08-08	131

In [16]: *# Convert the glucose values to float for calculations later*

```
new_cgm_data_df['glucose_value'] = new_cgm_data_df['glucose_value'].astype('float')  
  
# Check to make sure everything looks ok  
  
new_cgm_data_df.head()
```

Out[16]:

	date_entry	glucose_value
0	2019-08-08	140.0
1	2019-08-08	138.0
2	2019-08-08	136.0
3	2019-08-08	134.0
4	2019-08-08	131.0

```
In [18]: # In order to do aggregate calculations, set the index to the date
new_cgm_data_df = new_cgm_data_df.set_index('date_entry')

# Find the average glucose value by date

clean_cgm_df = new_cgm_data_df.groupby(new_cgm_data_df.index).mean()

# Display the data

clean_cgm_df.head()
```

Out[18]:

	glucose_value
date_entry	
2019-08-08	114.712803
2019-08-09	151.347222
2019-08-10	203.204861
2019-08-11	176.045139
2019-08-12	194.965278

```
In [19]: # Reset the index so that the data can be loaded to pgAdmin4

clean_cgm_df = clean_cgm_df.reset_index()

# Display the data

clean_cgm_df.head()
```

Out[19]:

	date_entry	glucose_value
0	2019-08-08	114.712803
1	2019-08-09	151.347222
2	2019-08-10	203.204861
3	2019-08-11	176.045139
4	2019-08-12	194.965278

```
In [21]: # Convert the sleep date time stamp to date only so that the date is in the same format as the glucose values table
# In order to join the two data frames, the date needs to be in the same format

new_sleep_data_df['End Time'] = pd.to_datetime(new_sleep_data_df['End Time']).dt.strftime('%Y-%m-%d')

# Display the data

new_sleep_data_df.head()
```

Out[21]:

	End Time	Minutes Asleep	Minutes Awake	Number of Awakenings	Time in Bed	Minutes REM Sleep	Minutes Light Sleep	Minutes Deep Sleep
0	2019-11-03	429	46	30	475	131.0	210.0	88.0
1	2019-11-02	192	31	19	223	29.0	135.0	28.0
2	2019-11-02	574	95	40	669	120.0	379.0	75.0
3	2019-11-01	247	45	16	292	50.0	161.0	36.0
4	2019-10-31	508	93	2	601	NaN	NaN	NaN

```
In [22]: # Rename the columns for clean data

new_sleep_data_df = new_sleep_data_df.rename(columns={'End Time': 'date_entry', 'Minutes Asleep': 'min_asleep', 'Minutes Awake': 'min_awake', 'Number of Awakenings': 'num_times_awake', 'Time in Bed': 'total_time_in_bed', 'Minutes REM Sleep': 'min_REM', 'Minutes Light Sleep': 'min_light', 'Minutes Deep Sleep': 'min_deep'})

# Display the data

new_sleep_data_df.head()
```

Out[22]:

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_light	min_deep
0	2019-11-03	429	46	30	475	131.0	210.0	88.0
1	2019-11-02	192	31	19	223	29.0	135.0	28.0
2	2019-11-02	574	95	40	669	120.0	379.0	75.0
3	2019-11-01	247	45	16	292	50.0	161.0	36.0
4	2019-10-31	508	93	2	601	NaN	NaN	NaN


```
In [24]: # Drop NaN values

new_sleep_data_df = new_sleep_data_df.dropna()

# Display the data

new_sleep_data_df.head()
```

```
Out[24]:
```

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_ligh
0	2019-11-03	429	46	30	475	131.0	210.0
1	2019-11-02	192	31	19	223	29.0	135.0
2	2019-11-02	574	95	40	669	120.0	379.0
3	2019-11-01	247	45	16	292	50.0	161.0
5	2019-10-30	221	31	16	252	39.0	138.0

```
In [25]: # Check the count of data

len(new_sleep_data_df)
```

```
Out[25]: 73
```

Load:

Create database and tables in pgAdmin4 and load the data above.

```
In [26]: # Create the connection to pgAdmin4

rds_connection_string = "postgres:*****@localhost:5432/T1D_db"
engine = create_engine(f'postgresql://{rds_connection_string}')
```

```
In [27]: # Check table names to make sure correct tables are there

engine.table_names()
```

```
Out[27]: ['cgm_data', 'sleep_data']
```

```
In [28]: # Load the data into the appropriate tables

clean_cgm_df.to_sql(name='cgm_data', con=engine, if_exists='replace', index=False)
new_sleep_data_df.to_sql(name='sleep_data', con=engine, if_exists='replace', index=False)
```

```
In [29]: # Read the data back to make sure data was loaded correctly

pd.read_sql_query('select * from cgm_data', con=engine).head()
```

Out[29]:

	date_entry	glucose_value
0	2019-08-08	114.712803
1	2019-08-09	151.347222
2	2019-08-10	203.204861
3	2019-08-11	176.045139
4	2019-08-12	194.965278

```
In [30]: # Read the data back to make sure data was loaded correctly

pd.read_sql_query('select * from sleep_data', con=engine).head()
```

Out[30]:

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_light
0	2019-11-03	429	46	30	475	131.0	210.0
1	2019-11-02	192	31	19	223	29.0	135.0
2	2019-11-02	574	95	40	669	120.0	379.0
3	2019-11-01	247	45	16	292	50.0	161.0
4	2019-10-30	221	31	16	252	39.0	138.0

Analysis:

Analysis on the clean data.

```
In [31]: # Join the two data frames for analysis
# Join on the date column

sleep_cgm_df = pd.merge(new_sleep_data_df, clean_cgm_df, on='date_entry', how=
'inner')

# Display data

sleep_cgm_df.head()
```

Out[31]:

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_ligh
0	2019-11-03	429	46	30	475	131.0	210.0
1	2019-11-02	192	31	19	223	29.0	135.0
2	2019-11-02	574	95	40	669	120.0	379.0
3	2019-11-01	247	45	16	292	50.0	161.0
4	2019-10-30	221	31	16	252	39.0	138.0

```
In [38]: # Create bins for the amount of sleep in minutes
# Each bin will have 120 minutes or 2 hours

bins = [0, 120, 240, 360, 480, 10000]

# Create labels for the bins
# Labels show the amount in hours

labels = ['<=2', '2<=4', '4<=6', '6<=8', '>8']

# Create a new column that shows where each value falls for the amount of sleep
p

sleep_cgm_df['hours'] = pd.cut(sleep_cgm_df['min_asleep'], bins=bins, labels=labels)

# Display the data

sleep_cgm_df.head()
```

```
Out[38]:
```

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_ligh
0	2019-11-03	429	46	30	475	131.0	210.0
1	2019-11-02	192	31	19	223	29.0	135.0
2	2019-11-02	574	95	40	669	120.0	379.0
3	2019-11-01	247	45	16	292	50.0	161.0
4	2019-10-30	221	31	16	252	39.0	138.0

```
In [42]: # Since data is in bins, find the average glucose value by amount of sleep

cgm_by_hours_sleep = pd.DataFrame(sleep_cgm_df.groupby(sleep_cgm_df['hours']).
mean()['glucose_value'])

# Display the data

cgm_by_hours_sleep.head()
```

```
Out[42]:
```

	glucose_value
hours	
<=2	NaN
2<=4	185.907280
4<=6	178.811851
6<=8	166.626095
>8	182.428819

In [43]: *# Find the max minutes of REM sleep in order to create appropriate bins*

```
sleep_cgm_df['min_REM'].max()
```

Out[43]: 174.0

In [44]: *# Create bins and labels for the REM sleep data*

```
REM_bins = [0, 60, 120, 180]
```

```
REM_labels = ['<=1', '1<=2', '>2']
```

Label the data into the appropriate bin

```
sleep_cgm_df['REM_hours'] = pd.cut(sleep_cgm_df['min_REM'], bins=REM_bins, labels=REM_labels)
```

Display the data

```
sleep_cgm_df.head()
```

Out[44]:

	date_entry	min_asleep	min_awake	num_times_awake	total_time_in_bed	min_REM	min_ligh
0	2019-11-03	429	46	30	475	131.0	210.0
1	2019-11-02	192	31	19	223	29.0	135.0
2	2019-11-02	574	95	40	669	120.0	379.0
3	2019-11-01	247	45	16	292	50.0	161.0
4	2019-10-30	221	31	16	252	39.0	138.0

In [45]: *# Find the average glucose value by REM sleep amount*

```
cgm_by_hours_sleep_2 = pd.DataFrame(sleep_cgm_df.groupby(sleep_cgm_df['REM_hours']).mean()['glucose_value'])
```

Display the data

```
cgm_by_hours_sleep_2
```

Out[45]:

	glucose_value
REM_hours	
<=1	179.996548
1<=2	173.700413
>2	175.453982

See PDF for further information